

# NYC Taxi Fare Prediction

## Group 1

Guilherme, Alex, Rakshit, Seonhye, Michael

# Agenda

1. Problem Background
2. Preprocessing
3. EDA
4. Models
5. GPU Acceleration
6. Results
7. Conclusion
8. Future Work
9. Utilized Class Topics

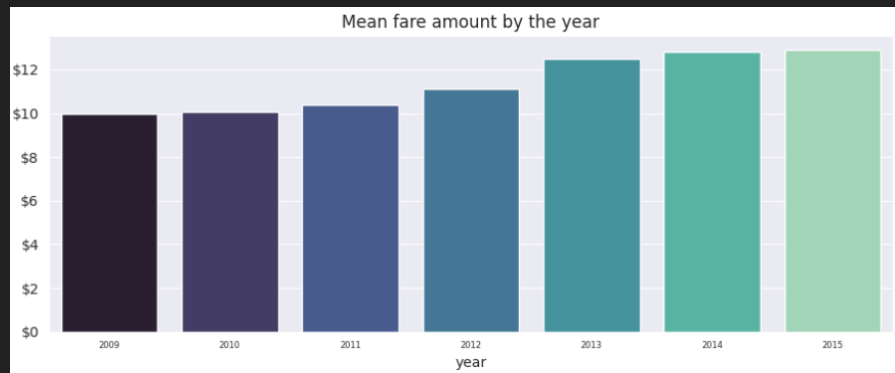
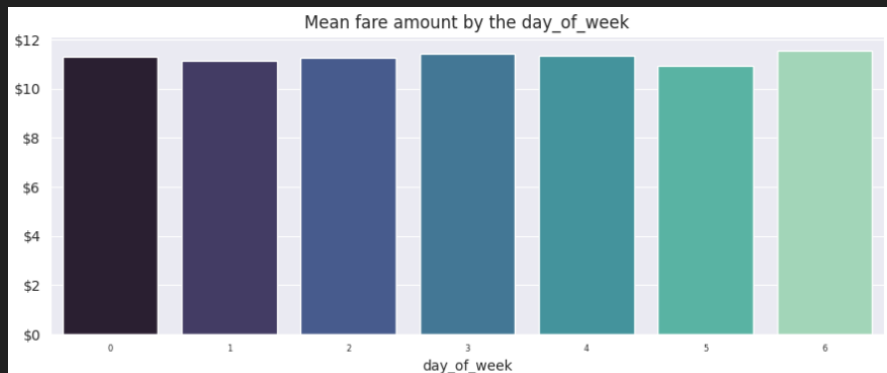
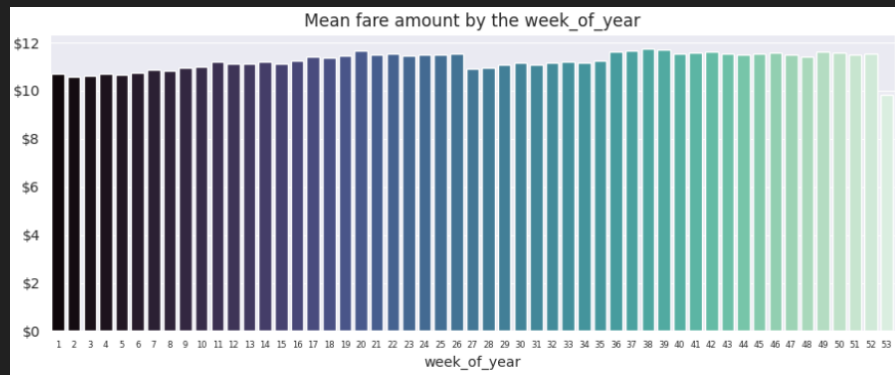
# Problem Background

- **Target:**
  - Predict fare of a taxi ride
- **Features:**
  - Pickup time and coordinates
  - Drop Off coordinates
  - Passenger count
- **Historical Results:**
  - \$3-5 MAE using just ride distance


# Preprocessing

- Drop rides that seem faulty/fake
  - 0 or more than 6 passengers
  - Below \$3 and above \$150
  - Low distance ( $< 1$  km) with excessively high fares ( $> \$40$ )
- Compute geographical distances
- Label pick-up and drop-off areas
  - Then one-hot encoded
- Dropped unused columns

# EDA - Features

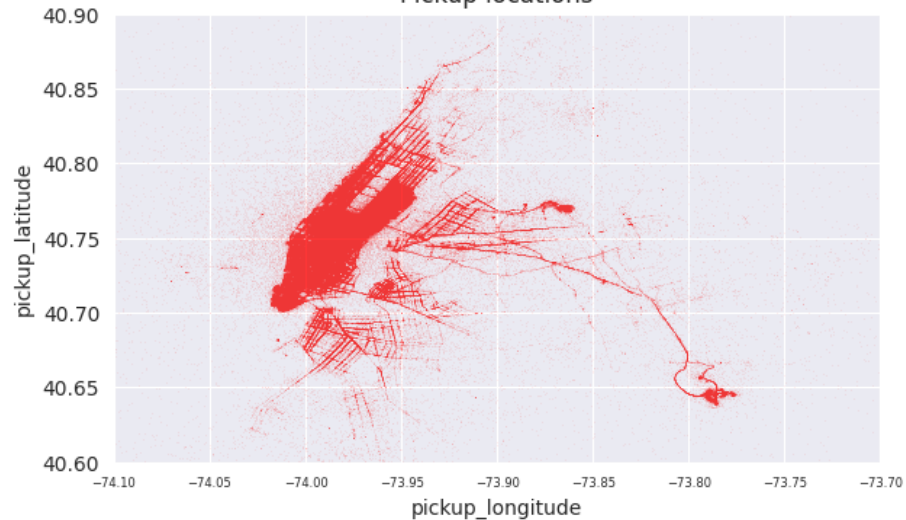


# EDA - Features




# EDA - Features


Pickup locations



Dropoff locations



# EDA - Features





# Models

- Linear Regression
    - Scikit-Learn and CuML versions
  - Random Forest
    - Scikit-Learn and CuML versions
  - XGBoost
- 
- ❖ Model Objective: Minimize RMSE of Taxi Fare prediction
    - RMSE penalizes larger errors
    - We report Mean Absolute Error (MAE) for human understanding

# GPU Acceleration

- We used CuDF and CuML libraries, developed by NVIDIA for work on NVIDIA GPUs.
- CuDF is a GPU-accelerated dataframe library that is API-compatible with Pandas
  - Utilizes similar Pandas syntax with GPU acceleration
- CuML is a machine learning library of machine learning algorithms optimized for GPU acceleration
  - Designed to work seamlessly with **CuDF**
- File size and device memory constrained how much data we could utilize
  - Required to use only 10% of given data to utilize **CuML**
  - Attempted PyTorch implementations failed entirely

# Results

## Device and Rows

**CPU;** 1 million rows

**CPU;** 5 million rows

**GPU;** 5 million rows

**CPU;** full data (55mil)

Models	train_mae (\$)	test_mae (\$)	train_r2	test_r2	train_time (s)
linear_regression	1.99	2.37	0.7742	0.7798	6.95
random_forest	1.7	1.92	0.8361	0.8379	521.22
xgboost	1.56	1.89	0.8683	0.8433	121.71

linear_regression	1.99	2.36	0.7737	0.7793	30.51
random_forest	1.71	1.92	0.8326	0.8377	3028.52
xgboost	1.63	1.86	0.8491	0.8471	135.56

linear_regression	2.05	2.41	0.7412	0.7624	1.91
random_forest	1.71	1.92	0.8308	0.8377	462.58
xgboost	1.63	1.86	0.8497	0.8468	9.94

linear_regression	1.99	2.36	0.7731	0.779	268.63
random_forest	1.72	1.92	0.8301	0.8373	41023.48
xgboost	1.63	1.85	0.8475	0.8481	1326.78

# Conclusions

- Best Model: XGBoost
  - Best MAEs and  $R^2$  scores
  - Ran very efficiently compared to Random Forest
- Linear Model close performance, but faster
  - More suitable for tasks demanding speed
- GPU fastest training speed
  - Suitable for tensor-like data
  - Harder to get devices with required memory
- Best Test MAE: \$1.85
  - Still large error for fair prediction, needs improvement
  - Insignificant MAE difference between 5mil and full, significant speed difference

# Future Work

- Utilize machines with more memory to use the full dataset
  - Properly tune model with such hardware
- Parallelize the computation more with more apt machines
- Utilize Google Maps API to try GPS distances and traffic levels as features
  - Required funding for practical use
- More precise EDA and cleaning
  - Finely remove outliers that cannot exist
    - Like being below base-rate
  - Breakdown pick-up and drop-off locations more finely
    - Like how going to an airport entails additional fees beyond regular rates

# Utilized Class Topics

- Shell
- Python Performance
- Optimization
- Parallel Programming
- Python for GPUs

*ANY QUESTIONS?*

